

高性能计算应用实践实验报告（lab6）

2024 秋 苏涵 2023311B25

实验环境

OS: Linux suh 5.15.153.1-microsoft-standard-WSL2 #1 SMP Fri Mar 29 23:14:13 UTC 2024

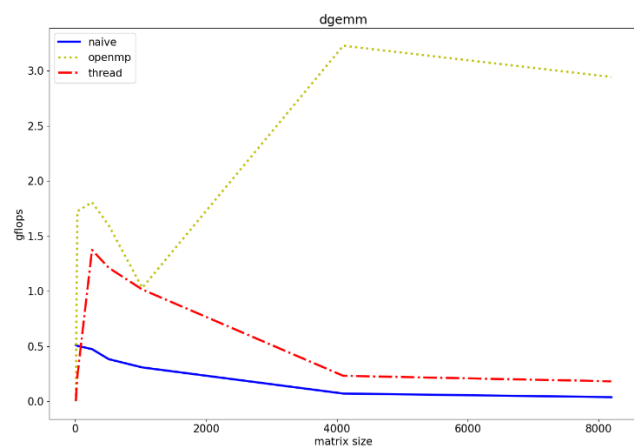

x86_64 x86_64 x86_64 GNU/Linux

Ubuntu 22.04.3 LTS

gcc: (Ubuntu 11.4.0-1ubuntu1~22.04) 11.4.0

CPU: 型号 11th Gen Intel(R) Core(TM) i7-11370H 频率 3.30GHz 核数 4

数据分析



当矩阵规模为 8×8 时, naive 的 gflops 大于其它方法, 其它情况下 gflops 从大到小排列依次为 cblas、openmp、thread、naive。且 cblas 的 gflops 远大于其他三种方法。

实现方式介绍与核心代码

naive

使用三重嵌套循环来计算矩阵乘法。外层两个循环遍历结果矩阵 C 的每个元素，内层循环计算对应元素的和

```
C naive_gemm.c > ...
1  #include "defs.h"
2
3  void MY_MMult(int m, int n, int k, double *a, int lda,
4               double *b, int ldb,
5               double *c, int ldc)
6  {
7      int i, j, p;
8
9      for (i = 0; i < m; i++) /* Loop over the rows of C */
10     {
11         for (j = 0; j < n; j++) /* Loop over the columns of C */
12         {
13             for (p = 0; p < k; p++)
14             {
15                 C(i, j) = C(i, j) + A(i, p) * B(p, j);
16             }
17         }
18     }
19 }
```

thread

将矩阵乘法任务分配到多个线程上，每个线程负责计算矩阵的一部分，最终合并结果

```
C pthreads_gemm.c > ...
1  #include "defs.h"
2  #include <pthread.h>
3
4  typedef struct {
5      int start_row;
6      int end_row;
7      double* A;
8      double* B;
9      double* C;
10     int m;
11     int n;
12 } ThreadData;
13
14
15 void* matrix_multiply(void* arg) {
16     ThreadData* data = (ThreadData*)arg;
17     int start_row = data->start_row;
18     int end_row = data->end_row;
19     double* A = data->A;
20     double* B = data->B;
21     double* C = data->C;
22     int n = data->n;
23
24     for (int i = start_row; i < end_row; i++) {
25         for (int j = 0; j < n; j++) {
26             for (int k = 0; k < n; k++) {
27                 C[i*n+j] = C[i*n+j] + A[i*n+k] * B[k*n+j];
28             }
29         }
30     }
31     pthread_exit(NULL);
32 }
33
34 void MY_MMult(int m, int n, int k, double *a, int lda,
35              double *b, int ldb,
36              double *c, int ldc)
37 {
38     int rows_per_thread = m / NUM_THREADS;
39     pthread_t threads[NUM_THREADS];
40     ThreadData thread_data[NUM_THREADS];
41
42     for (int i = 0; i < NUM_THREADS; i++) {
43         thread_data[i].start_row = i * rows_per_thread;
44         thread_data[i].end_row = (i + 1) * rows_per_thread;
45         thread_data[i].A = a;
46         thread_data[i].B = b;
47         thread_data[i].C = c;
48         thread_data[i].n = n;
49         pthread_create(&threads[i], NULL, matrix_multiply, (void*)&thread_data[i]);
50     }
51
52     for (int i = 0; i < NUM_THREADS; i++) {
53         pthread_join(threads[i], NULL);
54     }
55 }
```

cblas

使用 BLAS 库中的 `cblas_dgemm` 函数，自动处理矩阵乘法的计算

```
C cblas_gemm.c > ...
1  #include <cblas.h>
2  #include <stdio.h>
3
4  void MY_MMult(int m, int n, int k, double *a, int lda,
5              double *b, int ldb,
6              double *c, int ldc)
7  {
8      double alpha = 1.0;
9      double beta = 1.0;
10
11     cblas_dgemm(CblasRowMajor, CblasNoTrans, CblasNoTrans, m, n, k, alpha, a, lda, b, ldb, beta, c, ldc);
12
13 }
```

openmp

使用 OpenMP 并行化矩阵乘法，通过添加指令，利用多核处理器加速计算

```
C openmp_gemm.c > ...
1  #include "defs.h"
2  #include <omp.h>
3
4  void MY_MMult(int m, int n, int k, double *a, int lda,
5              double *b, int ldb,
6              double *c, int ldc)
7  {
8      #pragma omp parallel for
9      for (int i = 0; i < m; i++) /* Loop over the rows of C */
10     {
11         for (int j = 0; j < n; j++) /* Loop over the columns of C */
12         {
13             for (int p = 0; p < k; p++)
14             {
15                 c[i, j] = c[i, j] + A[i, p] * B(p, j);
16             }
17         }
18     }
19 }
```

lab3 问题

(1)

```
echo "date = `date`";" > $(DATA_DIR)/output_$(NEW).m
echo "version = 'naive_gemm';" >> $(DATA_DIR)/output_$(NEW).m
```

可通过查看 makefile 中 `new` 或运行时输出 `version` 判断调用的是哪个函数

```
24  $(BUILD_DIR)/test_MMult.x: $(OBJS) defs.h
25  $(LINKER) $(OBJS) $(LDFLAGS) -o $@

12  OBJS := $(BUILD_DIR)/util.o $(BUILD_DIR)/REF_MMult.o $(BUILD_DIR)/test_MMult.o $(BUILD_DIR)/$(NEW).o

2   NEW := naive_gemm
```

图中代码指定了使用 `new` 中的 `MY_MMult` 函数

(2)

```
33  @echo "date = `date`";" > $(DATA_DIR)/output_$(NEW).m
34  @echo "version = '$(NEW)';" >> $(DATA_DIR)/output_$(NEW).m
35  $(BUILD_DIR)/test_MMult.x >> $(DATA_DIR)/output_$(NEW).m
```

通过 makefile 中图代码，将运行后的数据写入到 `_data/output_MMult0.m`

thread 运行截图

```
wsl: A localhost proxy configuration was detected but not mirrored into WSL. WSL in NAT mode does not support lo
calhost proxies.
root@suh:~# top
top - 22:30:55 up 7:46, 1 user, load average: 4.16, 4.11, 4.44
Tasks: 71 total, 1 running, 67 sleeping, 3 stopped, 0 zombie
%Cpu(s): 48.7 us, 0.3 sy, 0.0 ni, 47.7 id, 0.0 wa, 0.0 hi, 3.3 si, 0.0 st
MiB Mem : 15914.0 total, 10750.5 free, 4403.4 used, 760.1 buff/cache
MiB Swap: 4096.0 total, 4096.0 free, 0.0 used, 11232.8 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
170190	root	20	0	2198456	2.0g	1788	S	400.0	12.9	128:52.10	test_MMult.x
1017	root	20	0	21.4g	279148	51300	S	2.0	1.7	14:27.65	node
1	root	20	0	165992	11352	8344	S	0.7	0.1	5:26.85	systemd
2	root	20	0	2476	1504	1384	S	0.0	0.0	0:00.03	init-systemd(Ub
6	root	20	0	2520	132	132	S	0.0	0.0	0:00.03	init
41	root	19	-1	31428	11464	10428	S	0.0	0.1	0:00.26	systemd-journal
65	root	20	0	22096	5888	4488	S	0.0	0.0	0:02.20	systemd-udev
76	root	20	0	153004	2204	4	S	0.0	0.0	0:00.00	snapfuse
78	root	20	0	377296	13268	272	S	0.0	0.1	0:01.07	snapfuse
80	root	20	0	153004	208	40	S	0.0	0.0	0:00.00	snapfuse
89	root	20	0	153136	2220	0	S	0.0	0.0	0:00.00	snapfuse
97	root	20	0	153004	2224	12	S	0.0	0.0	0:00.00	snapfuse
101	root	20	0	302532	7056	320	S	0.0	0.0	0:00.17	snapfuse
108	root	20	0	153004	212	52	S	0.0	0.0	0:00.00	snapfuse
111	root	20	0	302532	12976	264	S	0.0	0.1	0:01.65	snapfuse
120	systemd+	20	0	25540	12664	8472	S	0.0	0.1	0:00.21	systemd-resolve
136	root	20	0	4308	2744	2484	S	0.0	0.0	0:00.03	cron
137	message+	20	0	8584	4736	4196	S	0.0	0.0	0:00.06	dbus-daemon

```
root@suh:~# pstree
systemd--2*[agetty]
--cron
--dbus-daemon
--init-systemd(Ub)
|   SessionLeader--Relay(940)--4*[cpptools-srv--7*[{cpptools-srv}]]
|   |   sh--sh--sh--node--node--12*[{node}]
|   |   |   cpptools--17*[{cpptools}]+
|   |   |   |   node--6*[{node}]
|   |   |   |   |   pet
|   |   |   |   |   python3--python3
|   |   |   |   |   |   2*[{python3}]
|   |   |   |   |   |   11*[{node}]
|   |   |   |   |   |   node--bash
|   |   |   |   |   |   |   11*[{node}]
|   |   |   |   |   |   |   10*[{node}]
|   |   |   |   |   |   |   test_MMult.x
|   |   |   |   |   |   |   |   2*[test_MMult.x--7*[{test_MMult.x}]]
|   |   |   |   |   |   |   |   test_MMult.x--4*[{test_MMult.x}]
|   |   |   |   |   |   |   |   top
|   |   |   |   |   |   |   |   pstree
|   |   |   |   |   |   |   |   node--6*[{node}]
|   |   |   |   |   |   |   |   node--6*[{node}]
|   |   |   |   |   |   |   |   init--{init}
|   |   |   |   |   |   |   |   login--bash
|   |   |   |   |   |   |   |   {init-systemd(Ub)}
|   |   |   |   |   |   |   |   networkd-dispat
|   |   |   |   |   |   |   |   packagekitd--2*[{packagekitd}]
|   |   |   |   |   |   |   |   polkitd--2*[{polkitd}]
|   |   |   |   |   |   |   |   rsyslogd--3*[{rsyslogd}]
|   |   |   |   |   |   |   |   snapd--13*[{snapd}]
|   |   |   |   |   |   |   |   5*[snapfuse--2*[{snapfuse}]]
|   |   |   |   |   |   |   |   snapfuse--5*[{snapfuse}]
|   |   |   |   |   |   |   |   2*[snapfuse--4*[{snapfuse}]]
|   |   |   |   |   |   |   |   subiquity-serve--python3.10
|   |   |   |   |   |   |   |   |   python3
|   |   |   |   |   |   |   |   |   5*[{python3.10}]
|   |   |   |   |   |   |   |   systemd--(sd-pam)
|   |   |   |   |   |   |   |   systemd-journal
|   |   |   |   |   |   |   |   systemd-logind
|   |   |   |   |   |   |   |   systemd-resolve
|   |   |   |   |   |   |   |   systemd-udev
|   |   |   |   |   |   |   |   unattended-upgr--{unattended-upgr}
```

openmp 运行截图

```
top - 19:52:08 up 5:12, 1 user, load average: 0.74, 0.37, 0.44
Threads: 235 total, 9 running, 226 sleeping, 0 stopped, 0 zombie
%Cpu(s): 43.9 us, 0.5 sy, 0.0 ni, 55.2 id, 0.0 wa, 0.0 hi, 0.4 si, 0.0 st
MiB Mem : 15914.0 total, 14014.7 free, 1156.2 used, 743.2 buff/cache
MiB Swap: 4096.0 total, 4096.0 free, 0.0 used, 14480.7 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
126174	root	20	0	101440	43156	1932	R	99.9	0.3	0:11.70	test_MMult.x
126175	root	20	0	101440	43156	1932	R	36.0	0.3	0:01.36	test_MMult.x
126179	root	20	0	101440	43156	1932	R	36.0	0.3	0:01.36	test_MMult.x
126181	root	20	0	101440	43156	1932	R	36.0	0.3	0:01.36	test_MMult.x
126176	root	20	0	101440	43156	1932	R	35.7	0.3	0:01.35	test_MMult.x
126177	root	20	0	101440	43156	1932	R	35.7	0.3	0:01.35	test_MMult.x
126180	root	20	0	101440	43156	1932	R	35.7	0.3	0:01.35	test_MMult.x
126178	root	20	0	101440	43156	1932	R	35.3	0.3	0:01.33	test_MMult.x
1017	root	20	0	21.4g	279404	51300	S	1.0	1.7	10:42.89	node
1	root	20	0	165992	11352	8344	S	0.7	0.1	47:22.03	systemd
448	root	20	0	43376	37800	10216	S	0.3	0.2	6:07.23	python3
1025	root	20	0	21.4g	279404	51300	S	0.3	1.7	3:52.51	node
1028	root	20	0	21.4g	279404	51300	S	0.3	1.7	3:52.68	node
2063	root	20	0	1041020	62768	41236	S	0.3	0.4	0:19.83	node
33744	root	20	0	4259528	24940	11588	S	0.3	0.2	0:00.59	cpptools-srv
2	root	20	0	2476	1504	1384	S	0.0	0.0	0:00.05	init-systemd(Ub
8	root	20	0	2476	1504	1384	S	0.0	0.0	0:00.04	Interop
6	root	20	0	2520	132	132	S	0.0	0.0	0:00.00	init
7	root	20	0	2520	132	132	S	0.0	0.0	0:00.00	init
41	root	19	-1	31428	11448	10420	S	0.0	0.1	0:00.23	systemd-journal
65	root	20	0	22096	5888	4488	S	0.0	0.0	0:05.74	systemd-udev

```
systemd-2*[agetty]
-cron
-dbus-daemon
-init-systemd(Ub-SessionLeader-Relay(940)-4*[cpptools-srv-7*[{cpptools-srv}]]
-sh-sh-sh-node-node-12*[{node}]
-node-cpptools-17*[{cpptools}]+
-node-6*[{node}]
-pet
-python3-python3
-2*[{python3}]
-11*[{node}]
-node-bash
-11*[{node}]
-10*[{node}]
-SessionLeader-Relay(91608)-node-6*[{node}]
-SessionLeader-Relay(91618)-node-6*[{node}]
-SessionLeader-Relay(112933)-bash-test_MMult.x-7*[{test_MMult.x}]
-SessionLeader-Relay(114129)-bash-pstree
-SessionLeader-Relay(116605)-bash-top
-init-{init}
-login-bash
-{init-systemd(Ub}
-networkd-dispat
-packagekitd-2*[{packagekitd}]
-polkitd-2*[{polkitd}]
-rsyslogd-3*[{rsyslogd}]
-snapd-13*[{snapd}]
-5*[snapfuse-2*[{snapfuse}]]
-snapfuse-5*[{snapfuse}]
-2*[snapfuse-4*[{snapfuse}]]
-subiquity-serve-python3.10-python3
-5*[{python3.10}]
-systemd-(sd-pam)
-systemd-journal
-systemd-logind
-systemd-resolve
-systemd-udev
-unattended-upgr-{unattended-upgr}
```